

RAHUL MALIK

NLP Data Scientist

CONTACT

rahulmalik@email.com 
(123) 456-7890 
New York, NY 
[LinkedIn](#) 
[Github](#) 

EDUCATION

PhD
Natural Language
Processing (NLP)
University of Maryland
2019 - 2021
College Park, MD

B.S.
Statistics
Princeton University
2014 - 2018
Princeton, NJ

SKILLS

FAISS
Amazon SageMaker
LangGraph
MLflow
DeepEval
Apache Airflow
Instruction-design judgment
Human-in-the-loop
calibration

WORK EXPERIENCE

Data Scientist (NLP)

Amazon

January 2026 - current / New York, NY

- Reduced retrieval latency by 143 milliseconds with FAISS, sharpening answer relevance for customer-facing language experiences.
- Scaled model training and deployment in Amazon SageMaker, improving grounded-response precision by 6.8% across production NLP workflows.
- Analyzed and categorized multilingual intent patterns across 11 shopping-query classes, enhancing user search relevance and improving downstream response quality.
- **Generated \$1.84M in projected annual value** by surfacing failure modes that improved query understanding for high-volume search journeys.

Senior Machine Learning Engineer

Priceline

June 2023 - December 2025 / New York, NY

- Engineered itinerary-planning agents with LangGraph, **trimming handoff and orchestration overhead by 21 hours** per release cycle.
- Optimized experimentation with MLflow, increasing booking-assist conversion by 47 basis points for traveler support flows.
- Refactored semantic ranking features across 2.6M traveler utterances using NLP techniques such as feature engineering and model tuning, enhancing intent recognition accuracy by 7% in lodging and flight search.

Applied Scientist II

Microsoft

March 2022 - May 2023 / New York, NY

- Validated LLM response quality with DeepEval, **cutting regression detection time by 8 days** in pre-release evaluation pipelines.
- Automated retraining and scoring of 13 enterprise language models with Apache Airflow to ensure real-time model adjustments and minimize manual updates.
- Elevated task-completion accuracy by 7.3 points by isolating prompt and retrieval gaps in internal NLP product experiments.